

Corrective Synchronization via Trace Warping

Anonymous

1. Introduction

Concurrency control is a hard problem. While some thread interleavings are admissible (in particular, if they involve disjoint memory accesses), there are certain interleaving scenarios that must be inhibited to ensure serializability [?]. The goal is to automatically detect — with high precision and low overhead — the inadmissible interleavings, and ensure that that do not take effect.

Toward this end, there are currently two main synchronization paradigms:

- *Pessimistic synchronization*: In this approach, illegal interleaving scenarios are avoided conservatively by blocking the execution of one or more of the concurrent threads until the threat of incorrect execution has passed away. Locks, mutexes and semaphores are all examples of how to enforce mutual exclusion, or pessimistic synchronization.
- *Optimistic synchronization*: As an alternative to proactive, or pessimistic, synchronization, optimistic synchronization is essentially a reactive approach. The concurrency control system monitors execution, such that when an illegal interleaving scenario arises, it is detected as such and appropriate remediation steps are taken. A notable instance of this paradigm is transactional memory (TM) [?], where the system logs memory accesses by each of the threads, and is able to reverse the effects of a thread and abort/restart it.

Motivation The pessimistic approach is useful if critical sections are short, there is little available concurrency, and the involved memory locations are well known [?]. Optimistic synchronization is most effective when there is a high level of available concurrency. An example is graph algorithms, such as Boruvka, over graphs that are sparse and irregular [?].

Beyond these cases, however, there are many other situations of practical interest. As an illustrative example, we refer to the code fragment in Figure 1, extracted from the dyuproject project, where a shared Map object, (pointed-to by) `_convertors`, is manipulated by method `getConvertor()`.

Assume that different threads invoking this method are all attempting to simultaneously obtain the same `Convertor` object, which has not yet been created. Doing so optimistically would lead to multiple rollbacks, and thus poor performance. Mutual exclusion, on the other hand, would block all threads but one until the operation completes, which is far from optimal if `newConvertor()` is an expensive operation.

```
public Convertor getConvertor(
    Class cls, boolean create, boolean add) {
    Convertor convertor = _convertors.get(cls.getName());
    if (convertor == null && create) {
        convertor = newConvertor(cls, add);
        _convertors.putIfAbsent(cls.getName(), convertor);
    }
    return convertor;
}
```

Figure 1: Method `getConvertor()` from class `StandardConvertorCache` in project `dyuproject`

Our Approach We propose a novel synchronization paradigm, which is conceptually different from both the pessimistic and the optimistic approaches. In our approach, dubbed *corrective synchronization*, the correctness of multi-threaded execution is enforced after the fact, similarly to optimistic synchronization, though without rollbacks. Instead, the system automatically compensates, if necessary, for the effects of inadmissible interleavings by rewriting the program state as a transaction completes. This is done while accounting for the behavior of concurrent transactions, so as to guarantee serializability.

To illustrate our approach, we revisit the running example. Assume the following execution history:

| T_1 | T_2 |
|---|--|
| <code>_convertors.get()/null</code> | <code>_convertors.get()/null</code> |
| <code>if(...)</code> | <code>if(...)</code> |
| <code>newConvertor()/o₁</code> | <code>newConvertor()/o₂</code> |
| <code>_convertors.putIfAbsent()/null</code> | <code>_convertors.putIfAbsent()/o₁</code> |
| <code>return o₁</code> | <code>return o₂</code> |

This history is clearly nonserializable, as in any serializable history T_1 and T_2 would return the same `Convertor` instance. Correcting this execution involves the application of two actions to the exit state of T_2 . First, we point the local variable `convertor` to o_1 , rather than o_2 . Second, we fix the mapping under `_convertors` for key `cls.getName()` in the same way.

Note that the corrective actions above are of a general form, which is not limited to only two threads. For any number of threads, the corrected state would have one privileged thread deciding the return value (i.e., the value of `convertor`) for all threads, which would also be the value linked by the key under `_convertors`. Also note that the corrective actions are — relatively speaking — inexpensive, especially compared to the alternatives of either blocking or aborting/restarting all threads but one.

Two important challenges that we address in this paper w.r.t. the corrective synchronization paradigm are (i) how to compute correct poststates; and (ii) given an incorrect poststate, how to decide which correct poststate to transition to. We govern our discussion of these challenges by a formal framework, based on the push/pull model for transactions [?], with rigorous soundness guarantees. We also provide a clear statement of the limitations of corrective synchronization.


Beyond the formal details, this paper additionally addresses practical challenges, and in particular the question of how to implement corrective synchronization efficiently, such that it incurs low overhead. To this end, we present a solution based on static analysis to derive the correct poststates in relation to a given prestate. We have implemented a version of the analysis for shared Map data structures. Maps are used to represent the shared state of many Java programs [?], and so this first step toward a comprehensive static analysis for corrective synchronization is already of practical value.

say that the implementation does something very simple (assumes one transaction), yet illustrates something very powerful that is completely novel

Contributions This paper makes the following principal contributions:

1. **Corrective synchronization:** We present an alternative to both the pessimistic and the optimistic synchronization paradigms, dubbed *corrective synchronization*, whereby serializability is achieved neither via mutual exclusion nor via rollbacks, but through correction of the poststate according to a relational prestate/poststates specification.
2. **Formal guarantees:** We provide a formal description of corrective synchronization in terms of the push/pull model for transactions. This includes a correctness (or soundness) proof as well as a clear statement of limitations.
3. **Static analysis:** We have developed a static analysis to derive the prestate/poststates specification for programs that encode the shared state as one or more associative mappings. We describe the analysis in full formal detail.
4. **Implementation and evaluation:** We have created a prototype implementation of corrective synchronization assuming the shared state is represented as associative mappings. We discuss techniques and optimizations to achieve low overhead. We present experimental evidence in favor of corrective synchronization, where our subjects are derived from real-world Java applications.

2. Overview



3. Technical Background

In this section we describe a generic language of transactions and define an idealized semantics for concurrent transactions called the atomic semantics, in which there are no interleaved effects on the shared state. The model preliminaries generalize those provided previously [?]. We also define a notion of *good* configurations

and in the next section we will define how one can warp from a configuration that is not good to one that is.

Operations and States. We assume a set M of method calls or operations (e.g. `ht.put('a', 5)`). State is represented in terms of logs of operation records. An operation record (or, simply, an “operation”) $op = \langle m, \sigma_1, \sigma_2, id \rangle$ is a tuple consisting of the operation name m , a thread-local pre-stack σ_1 (method arguments), a thread-local post-stack σ_2 (method return values), and a unique identifier id . We assume a predicate $\text{fresh}(id)$ that holds provided that id is globally unique (details omitted for lack of space). In the atomic semantics defined below, the shared state ℓ : list op is an ordered list of operations. We use notations such as $\ell_1 \cdot \ell_2$ and $\ell \cdot op$ to mean append and appending a singleton, resp.

We require a prefix-closed predicate on operation lists allowed ℓ that indicates whether an operation log ℓ corresponds to a state. For convenience we will also write ℓ allows $\langle m, \sigma_1, \sigma_2, id \rangle$ which simply means allowed $\ell \cdot \langle m, \sigma_1, \sigma_2, id \rangle$. For example, if we have a simple TM based on memory read/write operations we expect allowed $\ell \cdot \{a := x, [x \mapsto 5], [x \mapsto 5, a \mapsto 5], id\}$, but \neg allowed $\ell \cdot \{a := x, [x \mapsto 5], [x \mapsto 5, a \mapsto 3], id\}$ or more elaborate specifications that involve multiple tasks. Ultimately, we expect the allowed predicate to be induced by the implementation’s operations on the state, $\llbracket op \rrbracket : \mathcal{P}(\text{State} \times \text{State})$, and initial states I .

We define a precongruence over operation logs $\ell_1 \preceq \ell_2$ coinductively, by requiring that all allowed extensions of the log ℓ_1 , are also allowed extension to the log ℓ_2 . We use a coinductive definition so that the precongruence can be defined up to all infinite suffixes.

$$\frac{\text{allowed } \ell_1 \Rightarrow \text{allowed } \ell_2 \quad \forall op. (\ell_1 \cdot op) \preceq (\ell_2 \cdot op)}{\ell_1 \preceq \ell_2}$$

We use a double-line here to indicate greatest fixpoint. Informally, the above definition says that there is no sequence of observations we can make of ℓ_2 , that we can’t also make of ℓ_1 . This is more general than just considering the set of states reached from executing the first log is included in the second: unobservable state differences are also permitted.

Language. Threads execute code c from some programming language that includes thread forking, transactions `tx c`, method names such as m , and a `skip` statement. As done elsewhere [?], we abstract away the programming language with a few semantic functions: **update this with pldi camera ready**

$c\Downarrow(m, c')$: Within a transaction, code c can be reduced to the pair (m, c') . That is, m is a next reachable method call in the reduction of c , with remaining code c' .

$c\downarrow(t, c')$: Outside of a transaction, code c can be reduced to the pair (t, c') . Here c' is the remaining code, and t is either a local state update, or a transaction or a thread fork.

fin(c): This predicate is true provided that there is a reduction of c to `skip` that does not encounter a method call.

These functions allow us to obtain a simple semantics, despite an expressive input language, by introducing functions to resolve nondeterminism between method operation names and at the end of a transaction. We assume that code is well-formed in that a single operation name m is always contained within a transaction.

Atomic Transition Systems. We next define a simple atomic semantics \xrightarrow{a} , in which transactions are executed instantly, without interruption from concurrent threads.

The \xrightarrow{a} rules AFIN, AFORK, ALOCAL shown in Figure 2(b) are similar to their counterparts in \xrightarrow{u} . However, the ATXN rule says that if thread executing code c_1 can reduce to a transaction

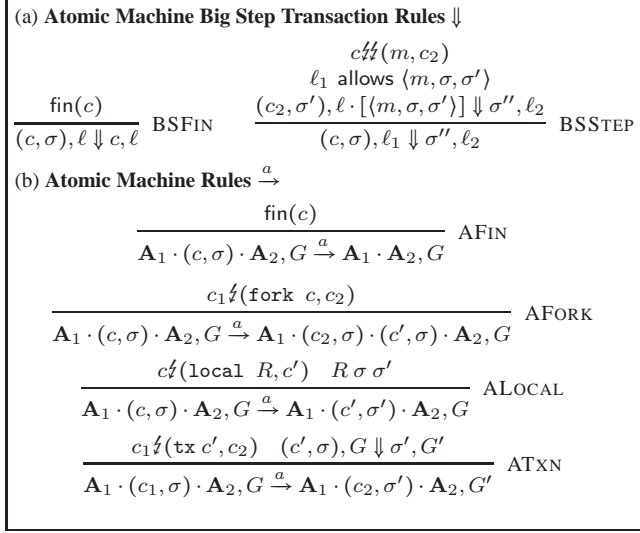


Figure 2: Atomic semantics of concurrent threads.

$\text{tx } c'$, then the transaction c' is executed atomically by the big step rules \Downarrow described next.

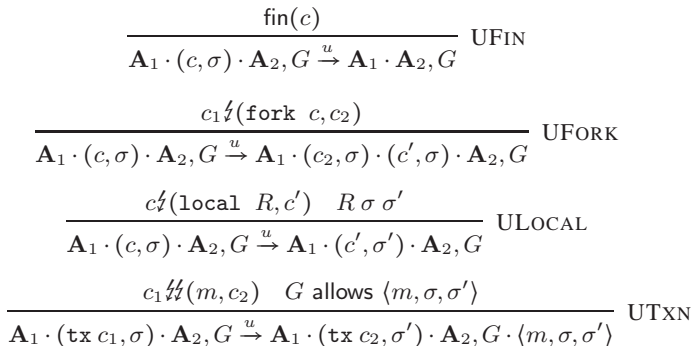
Figure 2(a) illustrates the big step semantics \Downarrow , which uses \Downarrow and $\text{fin}()$ (rules BSSSTEP and BSFIN, respectively). These rules scan through the nondeterminism in $\text{tx } c$ to find a next operation name m or a path to skip denoting the end of the transaction. BSSSTEP can be taken provided that the operation $\langle m, \sigma, \sigma' \rangle$ is permitted and that (c_2, σ') can be entirely reduced to (σ'', ℓ_2) .

4. Warping

The key idea of this paper is that when transactions run awry into an inconsistent state, rather than aborting them and starting again from the beginning, it may be possible to *correct* the shared/local state directly by directly modifying the state. Thus, the program continues as if it had not gone down the bad path to the inconsistent state.

We now formalize a simple version where, from an inconsistent state, the system warps to a state that was reachable in a serial interleaving. We later show this can be easily generalized to allow warping to more subtle states (such as those described by the Push/Pull model []). These *super-serial* (yet still serializable) allow for more possible destinations, making it more likely that a destination will be available when an inconsistent state is reached.

Unconstrained transition system. We begin with a generic transition system \xrightarrow{u} in which threads may interleave their effects however they please:



```

s ::= m.put(k, v)
      v = m.get(k)
      m.remove(k)
      v = m.putIfAbsent(k, v)
      v = new Value()
      v = null
      if(b) s1; else s2
      while(b) s1;
      s1; s2

b ::= x == NULL | m.containsKey(k) | !b1

```

Figure 3: Statements and conditions

The semantics is a relation \xrightarrow{u} over pairs consisting of a list of concurrent threads \mathbf{A} and a shared state ℓ . A single thread $(c, \sigma) \in \mathbf{A}$ is a code c and local state σ .

The unconstrained machine can take a UFIN step when there is a thread (c, σ) that can complete, *i.e.* $\text{fin}(c)$. The UFORK rule allows a new thread (c', σ) to be forked from thread (c, σ) . The ULOCAL rule involves manipulating the thread-local state σ to σ' . Finally, the UTXN rule allows a thread executing transaction code $\text{tx } c_1$ to take a single step to c_2 , applying the effects of m directly to the shared log G .

Corrective Warping and Committing. A simple version of warping (which we have implemented in Section ??) builds on an unconstrained transition system \xrightarrow{u} by adding a special WARPCMT rule. This rule attempts to perform a warp—replacing the current state with state that would have been reached in an atomic interleaving and then committing—or else, aborts. This is already more expressive than all existing notions of transactions which, in the face of an inconsistent state perform a (potentially partial) abort.

add the fact that each thread remembers the \mathbf{T}, G where they began

$$\frac{\text{fin}(c_1) \quad \text{Warp}(\mathbf{A}_1 \cdot (c_2, \sigma) \cdot \mathbf{A}_2, G, \mathbf{A}'_1 \cdot (c'_2, \sigma') \cdot \mathbf{A}'_2, G')}{\mathbf{A}_1 \cdot ((\text{tx } c_1, c_2), \sigma) \cdot \mathbf{A}_2, G \xrightarrow{u} \mathbf{A}'_1 \cdot (c'_2, \sigma') \cdot \mathbf{A}'_2, G'} \text{WARPCMT}$$

**how to jump to an atomic state? you haven't committed yet.
reference state issue.
conditions on warping - consistency.**

5. Language

As a proof of concept and preliminary practical study, we instantiate the theoretical framework we formalized in the Section 4 on the language in Figure 3, developing a static analysis to compute warping **targets**, and a dynamic system to warp.

The language is focused on a representative set of operations of the Java Map interface. In Figure 3, we represent by m the map shared among all the transactions, and k the shared key. The values inserted or read from the map might be a parameter of the transaction, or created through a new statement. Following the semantics of the Java library, our language supports (i) $v = m.get(k)$ that returns the value v related with key k , or null if k is not in the map, (ii) $m.remove(k)$ removes k from the map, (iii) $v = m.putIfAbsent(k, v)$ relates k to v in m if k is already in m and returns the previous value it was related to, (iv) $v = \text{new Value}(\dots)$ creates a new value, and (v) $v = \text{null}$ assigns null to variable v . In addition, our language support standard **if** and **while** statements, as well as concatenation of statements.

As Boolean conditions, the language supports checking if a variable is null, and if the map contains a key.

```

Value removeAttribute(Key k) {
  Value result = null;
  if (map.containsKey(k)) {
    result = m.get(k);
    m.remove(k);
  }
  return result;
}

boolean removeAttribute(Key k, Value v) {
  Value oldvalue = m.get(k);
  m.put(k, v);
  return oldvalue != null;
}

```

Figure 4: The running example inspired by class `ApplicationContext` of Apache Tomcat

5.1 Running Example

Figure 4 illustrates our running example. This code is inspired by **XXX**. The first type of transaction (`transaction1`) removes the value associated with the given key k , and returns it. Instead, the second type of transaction relates k with a given value v , and returns `true` if the key was already in the map. During the formalization of the static analysis and the warping system, we will refer to this running example where each transaction is instantiated multiple times, and all transactions conflict on the same key k .

6. Static Computation of Warp Destinations

6.1 Abstract Domain

Let `Var` and `HeapNode` be the set of variables and abstract heap nodes, respectively. We suppose that a special `null` value is part of `HeapNode`. Both keys and values are abstracted as heap nodes. As usual with heap abstractions, each heap node might represent one or many concrete nodes. Therefore, we suppose that a function `summary : HeapNode → {true, false}` is provided; `summary(h)` returns `true` if and only if h represents many concrete nodes (that is, it is a summary node). We define by $\text{Env} : \text{Var} \rightarrow \wp(\text{HeapNode})$ the set of (abstract) environments relating each variable to the set of heap nodes it might point to. A map is represented as a function $\text{Map} : \text{HeapNode} \rightarrow \wp(\text{HeapNode})$, connecting each key to the set of possible values it might be related to in the map. The value `null` represents that the key might not be in the map. For instance, $[n_1 \mapsto \{\text{null}, n_2\}]$ represents that the key n_1 might not be in the map, or it is in the map, and it is related to value n_2 . An abstract state is a pair made by an abstract environment and an abstract map. We augment this set with a special bottom value \perp to be used to represent that a statement is unreachable. Formally, $\Sigma = (\text{Env} \times \text{Map}) \cup \{\perp\}$.

The lattice structure is obtained by the point-wise application of set operators to elements in the codomain of abstract environments and functions. Therefore, the abstract lattice is defined as $\langle \Sigma, \subseteq, \dot{\cup} \rangle$, where \subseteq and $\dot{\cup}$ represents the point-wise application of set operators \subseteq and \cup , respectively.

Running example. Consider the first method in Figure 4. For instance, the abstract state $([k \mapsto \{n_1\}], [n_1 \mapsto \{\text{null}\}])$ represents that the key `key` is not in the map, while $([k \mapsto \{n_1\}], [n_1 \mapsto \{n_2\}])$ represents that it is in the map, and it is related to a value. Instead, for the second method, $([k \mapsto \{n_1\}, v \mapsto n_2], [n_1 \mapsto \{n_2\}])$ represents that k is in the map, and it is related to the value pointed by v .

6.2 Abstract Semantics

Figure 5 formalizes the abstract semantics of statements and Boolean conditions, that, given an abstract state (as defined in Section 6.1) and a statement or Boolean condition of the language introduced in Section 5, returns the abstract state resulting from the evaluation of the given statement on the given abstract state. We focus the formalization on abstract states in $\text{Env} \times \text{Map}$, since in case of \perp the abstract semantics always returns \perp itself.

`(put)` relates k to v in the map. In particular, if k points to a unique concrete node, it performs a so-called strong update, overwriting previous values related with k . Otherwise, it performs a weak update by adding to the previous values the new ones. `(get)` relates the assigned variable v to all the heap nodes of values that might be related with k in the map. Note that if k is not in the map, then the abstract map m relates it to a `null` node, and therefore this value is propagated to v then calling `get`, representing the concrete semantics of this statement. Similarly to `(put)`, `(rmv)` removes k from the map (by relating it to the singleton $\{\text{null}\}$) iff k points to a unique concrete node. Otherwise, it adds the heap node `null` to the heap nodes related to all the values pointed by k . `(pIA)` updates the map like `(put)` but only if the updated key node might have been absent, that is, when $\text{null} \in m(n)$. `new` creates a new heap node through `fresh(t)` (where t is the identifier of the transaction performing the creation), and assigns it to v . The number of nodes is kept bounded by parameterizing the analysis with an upper bound i such that (i) the first i nodes created by a transaction are all concrete nodes, and (ii) all the other nodes are represented by a summary node. Instead, `(null)` relates the given variable to the singleton $\{\text{null}\}$. Rules `(if)`, `(while)`, and `(cnc)` define the standard abstract semantics of `if`, `while`, and concatenation statements. The abstract semantics on Boolean conditions produces \perp statements if the given Boolean condition cannot hold on the given abstract semantics. Therefore, `(null)` returns \perp if the given variable x cannot be `null`, or a state relating x to the singleton $\{\text{null}\}$ otherwise. Vice-versa, `(!null)` returns \perp if x can be only `null`, or a state relating x to all its previous values except `null` otherwise. Similarly, `(cntK)` returns \perp if the given key k is surely not in the map, it refines the possible values of k if it is represented by a concrete node, or it simply returns the entry state otherwise. Vice-versa, `(!cntK)` returns \perp if k is surely in the map.

Running example. Consider again the first method in Figure 4. When we start from the abstract state $([k \mapsto \{n_1\}], [n_1 \mapsto \{\text{null}\}])$ (representing that k is not in the map), we obtain the abstract state $\sigma = ([k \mapsto \{n_1\}], [\text{result} \mapsto \{\text{null}\}], [n_1 \mapsto \{\text{null}\}])$ after the first statement by rule `(null)`. During the following computation of rule `(if)`, we consider:

1. when the Boolean condition `map.containsKey(k)` holds. When applying rule `(cntK)` on σ we obtain \perp since the node pointed by k is related to the singleton $\{\text{null}\}$ in the map, representing that the map does not contain the key k ; and
2. when `!map.containsKey(k)` holds. Rule `(!cntK)` applied to σ returns σ itself, since k is in relation only with `null` in the map.

`(if)` returns the upper bound of the two resulting states, that is $\perp \dot{\cup} \sigma = \sigma$, and the value pointed by `result` (that is, `null`) is returned. Therefore, our analysis computes on this example that, when the key is not in the map in the entry state, the method returns `null` and does not modify the map.

6.3 Over-approximating Serialized Executions

We now formalize how we compute an approximation of serialized executions relying on the abstract semantics defined in Section 6.2. Our system receives as input two types of transactions (e.g., the two

$$\begin{aligned}
\mathbb{S}[\mathbf{m.put}(k, v), (e, m)] &= \begin{cases} (e, m[n \mapsto e(v)]) & \text{if } e(k) = \{n\} \wedge \neg \text{summary}(n) \\ (e, m[n \mapsto m(n) \cup e(v) : n \in e(k)]) & \text{otherwise} \end{cases} & (\text{put}) \\
\mathbb{S}[v = \mathbf{m.get}(k), (e, m)] &= (e[v \mapsto \bigcup_{n \in e(k)} m(n)], m) & (\text{get}) \\
\mathbb{S}[\mathbf{m.remove}(k), (e, m)] &= \begin{cases} (e, m[n \mapsto \{\text{null}\}]) & \text{if } e(k) = \{n\} \wedge \neg \text{summary}(n) \\ (e, m[n \mapsto m(n) \cup \{\text{null}\} : n \in e(k)]) & \text{otherwise} \end{cases} & (\text{rmv}) \\
\mathbb{S}[v = \mathbf{m.putIfAbsent}(k, v), (e, m)] &= (\pi_1(\mathbb{S}[v = \mathbf{m.get}(k), (e, m)]), m') : \\
m' &= \begin{cases} (e, m[n \mapsto e(v)]) & \text{if } e(k) = \{n\} \wedge m(n) = \{\text{null}\} \\ (e, m[n \mapsto m(n) \cup e(v) : n \in e(k) \wedge \text{null} \in m(n)]) & \text{otherwise} \end{cases} & (\text{pIA}) \\
\mathbb{S}[v = \mathbf{new Value}(), (e, m)] &= (e[v \mapsto \text{fresh}(t)], m) & (\text{new}) \\
\mathbb{S}[v = \mathbf{new Value}(), (e, m)] &= (e[v \mapsto \{\text{null}\}], m) & (\text{null}) \\
\mathbb{S}[\mathbf{if}(b) s_1; \mathbf{else} s_2, (e, m)] &= \mathbb{S}[s_1, \mathbb{B}[b, (e, m)]] \dot{\cup} \mathbb{S}[s_2, \mathbb{B}[\neg b, (e, m)]] & (\text{if}) \\
\mathbb{S}[\mathbf{while}(b) s_1; , (e, m)] &= \mathbb{B}[\neg b, (e_1, m_1)] : (e_1, m_1) = \text{Lfp}_1^\varepsilon \lambda(e', m'). (e, m) \dot{\cup} \mathbb{S}[s_1, \mathbb{B}[b, (e', m')]] & (\text{while}) \\
\mathbb{S}[s_1; s_2, (e, m)] &= \mathbb{S}[s_2, \mathbb{S}[s_1, (e, m)]] & (\text{cnc}) \\
\mathbb{B}[x == \text{null}, (e, m)] &= \begin{cases} (e[x \mapsto \{\text{null}\}], m) & \text{if } \text{null} \in e(x) \\ \perp & \text{otherwise} \end{cases} & (\text{null}) \\
\mathbb{B}[\neg x == \text{null}, (e, m)] &= \begin{cases} (e[x \mapsto e(x) \setminus \{\text{null}\}], m) & \text{if } \exists n \in \text{HeapNode} : n \neq \text{null} \wedge n \in e(x) \\ \perp & \text{otherwise} \end{cases} & (\neg \text{null}) \\
\mathbb{B}[\mathbf{m.containsKey}(k), (e, m)] &= \begin{cases} \perp & \text{if } \forall n \in e(k) : m(n) = \{\text{null}\} \\ (e, m[n \mapsto m(n) \setminus \{\text{null}\}]) & \text{if } e(k) = \{n\} \wedge \text{summary}(n) \wedge m(n) \neq \{\text{null}\} \\ (e, m) & \text{otherwise} \end{cases} & (\text{cntK}) \\
\mathbb{B}[\neg \mathbf{m.containsKey}(k), (e, m)] &= \begin{cases} \perp & \text{if } \forall n \in e(k) : \text{null} \notin m(n) \\ (e, m[n \mapsto \{\text{null}\}]) & \text{if } e(k) = \{n\} \wedge \text{summary}(n) \wedge \text{null} \in m(n) \\ (e, m) & \text{otherwise} \end{cases} & (\neg \text{cntK})
\end{aligned}$$

Figure 5: Formal definition of the abstract semantics

methods in Figure 4), and returns a mapping from possible abstract entry states, to sets of possible exit states.

In particular, given two transactions t_1 and t_2 , we build up a control flow graph that represents possible serialized executions of many instances of the two transactions. In order to build up this serialized control flow graph, the local variables of transactions are rename. For instance, we can build up a serialized CFG where an instance of transaction t_1 (called t_{1_1}) is followed by a loop representing a (possible unbounded) sequence of executions of other instances of the same transaction represented by a loop of a *summary* instance of t_1 (called t_{1_n}), and this is followed by a similar structure for transaction t_2 (generating transactions t_{2_1} and t_{2_n}).

When applied to these serialized control flow graphs, the static analysis engine produces exit states that are possible results of serialized executions. Given two transactions, we represent by $\text{serializedCFGs}(t_1, t_2) = T$ the function that returns a set of serialized CFGs.

6.4 Extracting Possible Warping States

As we discussed in XXX, we need to compute warping that, given an entry state representing an observational equivalence class, is in an equivalence class that is reachable through a serialized execution. However, an abstract state in Σ might represent concrete states that are in different equivalence classes. For instance ($[k \mapsto$

$\{n_1\}], [n_1 \mapsto \{\text{null}, n_2\}])$ represents both that k is (if n_1 is related to n_2 in the abstract map) or is not (when n_1 is related to null). This abstract state therefore might concretize to states belonging to different equivalence classes, and it cannot used to define a warping destination.

Therefore, we define a predicate $\text{eqClass} : \Sigma \rightarrow \{\text{true}, \text{false}\}$ that, given an abstract state, holds iff it represents concrete states all in the same equivalence class. Formally,

$$\begin{aligned}
&\text{eqClass}(e, m) \\
&\quad \updownarrow \\
&\forall x \in \text{dom}(e) : |e(x)| = 1 \wedge e(x) = \{n_1\} \wedge \neg \text{summary}(n_1) \\
&\quad \forall n \in \text{dom}(m) : |m(n)| = 1 \wedge m(n) = \{n_2\} \wedge \neg \text{summary}(n_2)
\end{aligned}$$

Prove that if $\text{eqClass}(e, m)$ then all the concretized states from (e, m) are in the same equivalence class

Given two transactions t_1 and t_2 , we build up a set of possible entry states S such that $\forall (e, m) \in S : \text{eqClass}(e, m)$. We then compute the exit states for all the possible serialized CFGs and entry states, and we filter out only the ones that represents states in the same equivalence class. The results are represented as a function that relates each entry state to a set of possible exit states. This is formalized by the following warpDest function.

$$\text{warpDest}(\tau_1, \tau_2, S) = \{(e', m') : \\ \exists(e, m) \in S, \exists s \in \text{serializedCFGs}(\tau_1, \tau_2) : \\ (e', m') \in \mathbb{S}[\llbracket s, (e, m) \rrbracket] \wedge \text{eqClass}(e', m')\}$$

Not yet explained how we compute the possible entry states

6.5 Concrete state

Let $\Sigma = \gamma(\mathbf{T} \times G)$ be the concrete states, denoted $\sigma = (\bar{t}, g)$.

Discover a *Pietro* : $\hat{\Sigma} \rightarrow \hat{\Sigma} \rightarrow \wp(\hat{\Sigma})$, representing the current abstract state, the reference abstract state, and a set of possible destination abstract states.

7. Dynamic Warping

Given *Pietro*, the runtime system implements a function denoted *Peng* : $\sigma \rightarrow \hat{\sigma} \rightarrow \hat{\sigma} \rightarrow \sigma$.

Runtime tracks the current concrete state σ , current *abstract state* $\hat{\sigma}$ and the last *abstract reference state* $\hat{\sigma}_0$. Thus, we denote the runtime configuration as

$$c = \langle \sigma, \hat{\sigma}, \hat{\sigma}_0 \rangle \quad \text{or, expanding } \langle (\bar{t}, g), (\mathbf{T}, G), (\mathbf{T}_0, G_0) \rangle$$

That is, threads are in state \bar{t} , shared state g , tracked abstract state (\mathbf{T}, G) and tracked abstract reference state (\mathbf{T}_0, G_0) .

There are then the following rules for steps in the runtime system:

$$\begin{array}{c} \frac{\dots}{\langle \sigma, \hat{\sigma}, \hat{\sigma}_0 \rangle \hookrightarrow^* \langle \sigma', \hat{\sigma}', \hat{\sigma}_0 \rangle} \text{Diverge} \\[10pt] \frac{\hat{\sigma}' \in \text{Pietro}(\hat{\sigma}_0, \hat{\sigma}) \quad \sigma' = \text{Peng}(\sigma, \hat{\sigma}, \hat{\sigma}')}{\langle \sigma, \hat{\sigma}, \hat{\sigma}_0 \rangle \hookrightarrow \langle \sigma', \hat{\sigma}', \hat{\sigma}_0 \rangle} \text{Warp} \\[10pt] \frac{\dots}{\langle \sigma, \hat{\sigma}, \hat{\sigma}_0 \rangle \hookrightarrow \langle \sigma', \hat{\sigma}', \hat{\sigma}' \rangle} \text{Commit} \\[10pt] \frac{\text{fix } g \in \gamma(G) \quad \mathbf{T}, G \xrightarrow{PP} \mathbf{T}', G' \quad g' \in \gamma(G')}{\langle \sigma, \hat{\sigma}, \hat{\sigma}_0 \rangle \hookrightarrow \langle \sigma', \hat{\sigma}', \hat{\sigma}_0 \rangle} \text{Step} \end{array}$$

Pietro ensures that $\hat{\sigma}'$ is reachable from $\hat{\sigma}_0$.

Peng ensures that $\sigma \in \gamma\hat{\sigma}$ and that you always warp before you commit (or you always eventually warp)

8. Experimental results

8.1 Implementation

Execution Summary The static analysis phase computes a set of execution summaries, each representing a legal execution, which are used as the input of the dynamic analysis phase. Each execution summary describes (i) the return value of each transaction instance and (ii) the final map state

In our experiments, we have two threads running two types of transactions (TX_1 and TX_2) respectively. One exemplary execution summary is, $[key \rightarrow v_1^1, r_1^1 : v_1^1, r_1^n : v_1^1, r_2^1 : v_1^1, r_2^n : v_1^1]$, where r_1^1 is the symbolic form of the return value of TX_1 (1st instance of TX_1), v_1^1 is the symbolic form of the value put by the transaction instance, and $r_1^1 : v_1^1$ describes the return value of the transaction instance. The readers may notice r_1^n . This symbol represents the return value of the instances that are not explicitly specified. Here n is determined by the capability of the static analysis. At runtime, we track the instance order and use n by default if the number is not specified in the summary. The tracking of the instance order requires us to synchronize the first $n - 1$ instances.

However, the n is typically a small number, and the synchronization overhead is negligible. We do not need to synchronize starting from the n th instance because all these instances are represented as n by default. In addition, $key \rightarrow v_1^1$ tells what the key is associated with in the final map state.

The execution summary is not limited to the above basic case. In general, it is initial-state sensitive, schedule-oblivious and site-sensitive. First, the initial state about whether the key is mapped to some value affects the computation of the execution summaries. Therefore, we prefix each execution summary with a specific initial state, e.g., the state with initial key $[key^{init} \rightarrow v^{init}]$ or the state without initial key $[key^{init} \rightarrow void]$. Second, our execution summary is designed to be oblivious of the schedules. The schedule-obliviousness frees us from tracking the schedules at runtime and avoids the high tracking overhead. Third, the execution summary is site-sensitive. A transaction TX_1^1 may put a value dynamically created at site A to the map. We symbolically represent the value using the site and the occurrence of the site (inside the transaction instance), e.g., $A1$ in TX_1^1 . The extent to which we can distinguish the occurrences is determined by the static analysis, e.g., how many iterations can be unrolled.

Runtime System The runtime works as follows. We first use the counters to track the transaction instance and assemble the symbolic return accordingly, e.g., r_1^1 . Then we search for the symbolic value in the execution summaries, e.g., v_1^1 . Last, using the symbolic value as a key, we look up the cache, which is maintained to associate each symbolic value with a runtime value, for the concrete runtime value.

The second step, i.e., searching for symbolic value in the summaries, is challenging. The searching is demanded on the fly at the return of each transaction instance. However, the returns should be consistent with each other such that the warped execution represents a realistic execution. In the other words, the return values should be searched for in the same execution summary. To achieve this, we implement an on-the-fly pruning algorithm. At the first return, it finds the symbolic value in a randomly picked execution summary. After the value is picked, the execution summaries with different value for the return will be pruned, leading to a smaller solution space. The algorithm is iteratively applied. In addition, to achieve initial-state sensitivity, we also reduce the solution space based on the initial states.

Another tricky issue is, at some return, the symbolic value found in the execution summaries may have not been associated with any concrete runtime value yet. For example, at the return of TX_2^1 , we find the returned value as v_1^1 but TX_1^1 has not been executed, then we cannot find the runtime value associated with v_1^1 in the cache. This case happens because the execution summary is computed for certain schedule, which differs from the current schedule. In this case, we apply the notify/wait primitives to synchronize the cache lookup and cache maintenance. Even more complex, the returned value may never be put into the cache if the value creation site is disabled in the current schedule, e.g., its guarding branch condition is false. We simply remove all the guarding branches for the creation site such that it is always executed. This simple strategy is based on the insight that we treat the transaction as a blackbox and we only care about its return values. It is unnecessary to preserve the internal program semantics.

We also implemented the optimization for the caching. From the summaries, we see that only the values put by the first few transactions are used. Therefore, we only record such values into the cache and discard the rest at runtime.

8.2 Evaluation

For performance measurement, we compare 4 versions with different types of synchronizations, the original version *orig*, the ab-